

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score
df = pd.read_csv('diabetes.csv')
df.head()

```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome			
0	6	148	72	35	0	33.6		0.627	50	1		
1	1	85	66	29	0	26.6		0.351	31	0		
2	8	183	64	0	0	23.3		0.672	32	1		
3	1	89	66	23	94	28.1		0.167	21	0		
4	0	137	40	35	168	43.1		2.288	33	1		

Next steps: [Generate code with df](#) [New interactive sheet](#)

```

# 1. Replace zero values in key features with column median
key_features = ['Glucose','BloodPressure','SkinThickness','Insulin','BMI']

for col in key_features:
    median = df.loc[df[col] != 0, col].median()
    df.loc[df[col] == 0, col] = median

# 2. Replace the first row's glucose value with the maximum glucose
df.loc[0, 'Glucose'] = df['Glucose'].max()

# 3. For records with the lowest age, replace glucose values with the minimum
min_age = df['Age'].min()
df.loc[df['Age'] == min_age, 'Glucose'] = df['Glucose'].min()

df.head()

```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome			
0	6	199	72	35	125	33.6		0.627	50	1		
1	1	85	66	29	125	26.6		0.351	31	0		
2	8	183	64	29	125	23.3		0.672	32	1		
3	1	44	66	23	94	28.1		0.167	21	0		
4	0	137	40	35	168	43.1		2.288	33	1		

Next steps: [Generate code with df](#) [New interactive sheet](#)

```

print(df.info())
print(df.describe())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Pregnancies      768 non-null    int64  
 1   Glucose          768 non-null    int64  
 2   BloodPressure    768 non-null    int64  
 3   SkinThickness    768 non-null    int64  
 4   Insulin          768 non-null    int64  
 5   BMI              768 non-null    float64 
 6   DiabetesPedigreeFunction 768 non-null    float64 
 7   Age              768 non-null    int64  
 8   Outcome          768 non-null    int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
None
Pregnancies      Glucose      BloodPressure  SkinThickness  Insulin  \
count    768.000000  768.000000  768.000000  768.000000  768.000000

```

```
mean    3.845052  116.294271   72.386719  29.108073  140.671875  
std     3.369578  36.797403   12.096642  8.791221  86.383060  
min     0.000000  44.000000   24.000000  7.000000  14.000000  
25%    1.000000  95.000000   64.000000  25.000000  121.500000  
50%    3.000000  115.000000   72.000000  29.000000  125.000000  
75%    6.000000  140.000000   80.000000  32.000000  127.250000  
max    17.000000  199.000000  122.000000  99.000000  846.000000
```

```
BMI  DiabetesPedigreeFunction      Age      Outcome  
count 768.000000                768.000000  768.000000  768.000000  
mean  32.455208                0.471876   33.240885  0.348958  
std   6.875177                0.331329   11.760232  0.476951  
min   18.200000                0.078000   21.000000  0.000000  
25%   27.500000                0.243750   24.000000  0.000000  
50%   32.300000                0.372500   29.000000  0.000000  
75%   36.600000                0.626250   41.000000  1.000000  
max   67.100000                2.420000   81.000000  1.000000
```

```
X = df.drop('Outcome', axis=1)  
y = df['Outcome']  
  
scaler = StandardScaler()  
X_scaled = scaler.fit_transform(X)  
  
X_train, X_test, y_train, y_test = train_test_split(  
    X_scaled, y, test_size=0.2, random_state=42  
)
```

```
model = LinearRegression()  
model.fit(X_train, y_train)
```

```
LinearRegression
```

```
y_pred_cont = model.predict(X_test)  
y_pred = np.round(y_pred_cont).astype(int)  
y_pred = np.clip(y_pred, 0, 1)
```

```
acc = accuracy_score(y_test, y_pred)  
prec = precision_score(y_test, y_pred, zero_division=0)  
rec = recall_score(y_test, y_pred, zero_division=0)  
f1 = f1_score(y_test, y_pred, zero_division=0)  
cm = confusion_matrix(y_test, y_pred)  
  
print("Accuracy:", acc)  
print("Precision:", prec)  
print("Recall:", rec)  
print("F1-score:", f1)  
print("Confusion Matrix:\n", cm)
```

```
Accuracy: 0.7597402597402597  
Precision: 0.68  
Recall: 0.6181818181818182  
F1-score: 0.6476190476190476  
Confusion Matrix:  
[[83 16]  
 [21 34]]
```

```
# 1. Calculate and print the accuracy score  
accuracy = accuracy_score(y_test, y_pred_binary)  
print(f"Accuracy: {accuracy:.4f}")  
  
# 2. Calculate the confusion matrix  
conf_matrix = confusion_matrix(y_test, y_pred_binary)  
  
# 3. Create a heatmap of the confusion matrix  
plt.figure(figsize=(8, 6))  
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',  
            xticklabels=['Predicted 0', 'Predicted 1'],  
            yticklabels=['Actual 0', 'Actual 1'])  
plt.xlabel('Predicted')
```


```

plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

# 4. Print a classification report
class_report = classification_report(y_test, y_pred_binary)
print("\nClassification Report:")
print(class_report)

```

Accuracy: 0.7532



```

Classification Report:
      precision    recall  f1-score   support

          0       0.79      0.84      0.81      99
          1       0.67      0.60      0.63      55

   accuracy                           0.75      154
  macro avg       0.73      0.72      0.72      154
weighted avg       0.75      0.75      0.75      154

```